



“SMS Detection using NLP Techniques”



- BENIGHENIMA Hafsa
- GHANDOUZ Amina

Agenda

- * Introduction
- * Dataset overview
- * Baseline approach
- * Transformer-based Approach

Introduction

Every day, millions of SMS messages are sent – many of them are spam, fraudulent, or malicious.

- These messages try to scam users, spread malware, or waste user time.
- Traditional keyword-based filters are ineffective, easily bypassed, and not adaptive.

We propose an NLP-based SMS spam detection system that:


- Uses Natural Language Processing (NLP) to understand message content.
- Applies machine learning and deep learning models to classify messages as spam or ham (not spam).

Dataset Overview

- Source: Kaggle
- Number of rows: 5,574 messages
- Classes: Spam / Ham

	Label	text
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

- Class Distribution



Baseline Approach

1. Data Preprocessing Pipeline

1.1 Train-Test Split

- Split dataset: 80% training / 20% testing
- Used stratified splitting to maintain class balance

1.2. Text Normalization

- Lowercasing
- Removing special characters

1.3. Text Vectorization

- Transformed text using TF-IDF
- Applied only to training data to avoid data leakage and ensure test set remains unseen and evaluation is fair

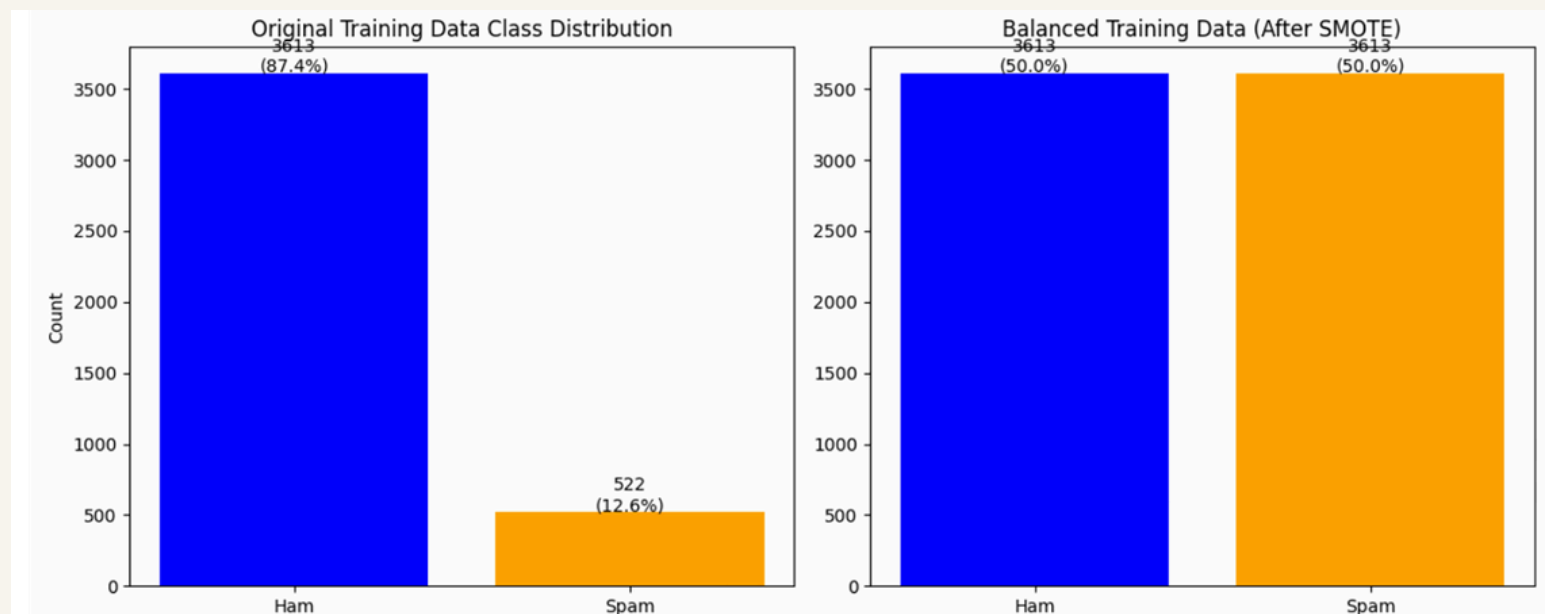
What is data leakage?

- TF: how often the word appears in a document.
- IDF: how rare the word is across all documents.

If we fit TF-IDF on the entire dataset (training + test), then the IDF part includes information from the test data, which is supposed to be unseen during training. This creates a data leakage problem – the model indirectly “sees” part of the test set, leading to unrealistically high performance during evaluation.

1.4 Addressing Class Imbalance

- Applied SMOTE after tf-idf to oversample minority class in training set

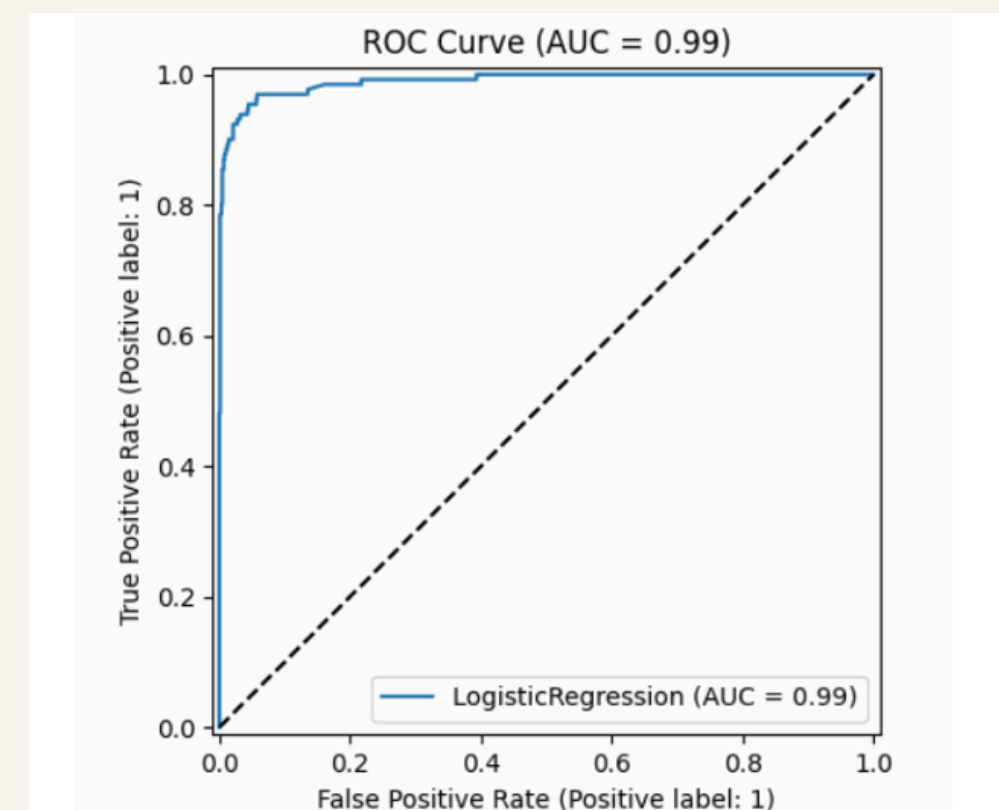



2. Model Training Process

- Logistic Regression (LR)
- Gaussian Naive Bayes (GNB)
- Random Forest (RF)

Model	Recall	F1-Score	ROC-AUC
Logistic Regression	0.98	0.98	0.99
Gaussian Naive Bayes	0.88	0.89	0.87
Random Forest	0.95	0.95	0.97

After evaluating all models using comprehensive classification metrics – including accuracy, precision, recall, F1-score, confusion matrix, and ROC curves – we observed that Logistic Regression consistently outperformed the others.




Transformer-based Approach



Why Transformers for SMS Spam Detection?

- Traditional models rely on handcrafted features (like TF-IDF).
- Transformers (like BERT) understand context, semantics, and word relationships.

Transformer-based Approach



Model Used & Setup

- Model: BERT
- Library/Framework: Hugging Face Transformers and PyTorch

BERT Input

- Input: A sentence is tokenized using WordPiece tokenizer
- Special tokens added:
- [CLS] = Classification token added at the start of every input.
- [SEP] = Separator token added at the end (or between sentences if there are two).

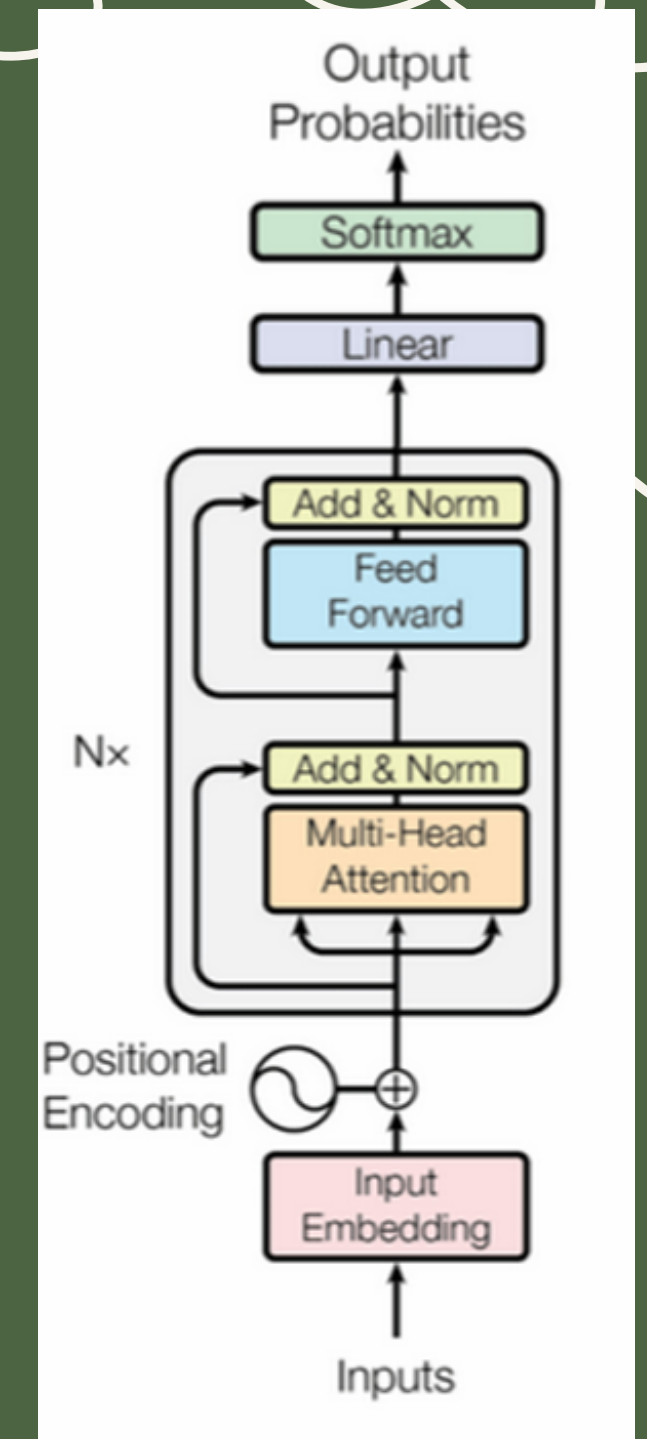
BERT Architecture?

we used BERT Base, that have:

- 12 transformer encoder layers and self-attention heads
- Add a final classification layer with 2 neurons, corresponding to the Ham and Spam classes.
- 768-dimensional hidden representations
- A total of around 110 million trainable parameters

How BERT Works?

- Each token is converted into an embedding vector.
- The whole sequence passes through multiple Transformer layers.
- The output embedding corresponding to [CLS] is used by a classifier to predict: Spam or Ham



Data preprocessing

- Used a cleaned SMS dataset (no missing values or duplicates)
- Split into 80% training / 20% test
- Applied stratified sampling to preserve spam/ham ratio across both sets

Tokenization and Input Formatting

- Used bert-base-uncased tokenizer (based on WordPiece)
- Automatically lowercases text
- Converts SMS into token IDs
- Padded/truncated to 128 tokens
- Returns:
 - input_ids (tokens)
 - attention_mask (1 = real token, 0 = padding)

Addressing Class Imbalance

- Computed class weights based on inverse class frequency

$$\text{weight_class} = \text{total_samples} / (\text{num_classes} \times \text{samples_in_class})$$

- Integrated weights into the cross-entropy loss
- Spam (minority class) received higher weight
- This helps model better detect spam by penalizing misclassifications more heavily

Model Configuration and Training

- Loss function: CrossEntropyLoss with class weights
- Optimizer: Adam with learning rate $2e-5$
- Epochs: 3
- Batch size: 16

Evaluation


Accuracy: 0.9923

Coss : 0.0235

Classification report:

Classification Report:				
	precision	recall	f1-score	support
Ham	0.99	1.00	1.00	903
Spam	0.99	0.95	0.97	131
accuracy			0.99	1034
macro avg	0.99	0.97	0.98	1034
weighted avg	0.99	0.99	0.99	1034

Confusion matrix:



The image features a dark green background with the text 'THANK YOU!' in a gold-colored serif font. In the top right corner, there are two overlapping swirls, one in white and one in gold. In the bottom left corner, there are also two overlapping swirls, one in white and one in gold.

THANK YOU!